

Assignment On Categorical Data Analysis

Name: Md Sadiq Doula Rahi

ID: 12110072

Department of Statistics

Begum Rokeya University, Rangpur

1. Consider the lung cancer data and solve the following questions using Python.

Basic EDA Questions i) Summarize the distribution of Age. Is it normally distributed or skewed? ii) What proportion of individuals are smokers vs non-smokers? iii) Which income group has the highest frequency? iv) What percentage of individuals are exposed to high pollution? v) What is the overall prevalence of lung disease?

Association & Crosstab i) Is there an association between Smoking and Lung Disease? ii) Interpret using a contingency table. iii) Does Pollution level affect Lung Disease? iv) Compare lung disease prevalence across Income groups. v) Which factor appears to have the strongest relationship with lung disease?

Statistical Testing i) Perform a Chi-square test for: Smoking vs Lung Disease and Pollution vs Lung Disease. ii) When should you use Fisher's Exact Test instead of Chi-square? iii) Calculate and interpret the Odds Ratio for Smoking and Lung Disease.

Correlation & Interpretation i) Interpret the correlation between: Smoking and Lung Disease, Age and Lung Disease. ii) Why is correlation not sufficient for causal inference?

Logistic Regression i) Fit a logistic regression model: Lung Disease ~ Smoking + Age + Pollution + Income. ii) Which variables are statistically significant predictors? ii) Interpret the odds ratio of Smoking. iii) What happens to the effect of Smoking after adjusting for Pollution?

Advanced Modeling i) Add an interaction term (Smoking x Pollution). ii) Is the interaction significant? iii) Interpret the interaction: Does pollution amplify smoking risk?

Research question: using the lung disease dataset, perform exploratory data analysis and fit a logistic regression model to identify key determinants of lung disease. interpret your finds and discuss public health implications. now give solution using python

```
[10]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from scipy.stats import chi2_contingency, fisher_exact

import statsmodels.api as sm

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import (
    confusion_matrix,
    classification_report,
    roc_curve,
    roc_auc_score
)

# 2. LOAD DATASET
df = pd.read_csv("lung_disease.csv")
```

```

print("\nFIRST 5 ROWS")
print(df.head())

print("\nCOLUMN NAMES")
print(df.columns)

print("\nDATA INFO")
print(df.info())

# 3. BASIC EDA

# AGE DISTRIBUTION

print("\nAGE SUMMARY")
print(df['Age'].describe())

sns.histplot(df['Age'], kde=True, color='skyblue')

plt.title("Age Distribution")
plt.xlabel("Age")
plt.ylabel("Frequency")

plt.show()

# SMOKERS VS NON-SMOKERS

print("\nSMOKING COUNTS")
print(df['Smoking'].value_counts())

sns.countplot(x='Smoking', data=df)

plt.title("Smoking Distribution")

plt.show()

# INCOME GROUP FREQUENCY

print("\nINCOME GROUP COUNTS")

income_counts = df['Income'].value_counts()

print(income_counts)

print("\nMOST COMMON INCOME GROUP")
print(income_counts.idxmax())

# HIGH POLLUTION PERCENTAGE

```

```

high_pollution_percentage = (
    (df['Pollution'] == 'High').mean() * 100
)

print(f"\nHIGH POLLUTION EXPOSURE: {high_pollution_percentage:.2f}%")

# LUNG DISEASE PREVALENCE

lung_prevalence = (
    (df['LungDisease'] == 1).mean() * 100
)

print(f"\nLUNG DISEASE PREVALENCE: {lung_prevalence:.2f}%")

# 4. CROSSTAB & ASSOCIATION

# SMOKING VS LUNG DISEASE

crosstab_smoking = pd.crosstab(
    df['Smoking'],
    df['LungDisease']
)

print("\nCROSSTAB: SMOKING VS LUNG DISEASE")
print(crosstab_smoking)

sns.heatmap(
    crosstab_smoking,
    annot=True,
    fmt='d',
    cmap='Blues'
)

plt.title("Smoking vs Lung Disease")

plt.show()

# 5. POLLUTION EFFECT

income_disease = pd.crosstab(
    df['Income'],
    df['LungDisease'],
    normalize='index'
) * 100

print("\nLUNG DISEASE BY INCOME GROUP (%)")

```

```

print(income_disease)

income_disease.plot(kind='bar', stacked=True)

plt.title("Lung Disease Across Income Groups")
plt.ylabel("Percentage")

plt.show()

# 6. STATISTICAL TESTING

# CHI-SQUARE TEST

chi2, p, dof, expected = chi2_contingency(crosstab_smoking)

print("\nCHI-SQUARE TEST")

print("Chi-square statistic:", chi2)
print("P-value:", p)

if p < 0.05:
    print("Significant association exists")
else:
    print("No significant association")

# POLLUTION VS LUNG DISEASE

crosstab_pollution = pd.crosstab(
    df['Pollution'],
    df['LungDisease']
)

chi2_p, p_p, dof_p, expected_p = chi2_contingency(
    crosstab_pollution
)

print("\nPOLLUTION VS LUNG DISEASE")

print("Chi-square statistic:", chi2_p)
print("P-value:", p_p)

# FISHER'S EXACT TEST

print("""

```

```

Use Fisher's Exact Test when:
- Sample size is small
- Expected counts are low
- Data is 2x2
"""

if crosstab_smoking.shape == (2, 2):

    oddsratio, fisher_p = fisher_exact(crosstab_smoking)

    print("\nFISHER'S EXACT TEST")

    print("Odds Ratio:", oddsratio)
    print("P-value:", fisher_p)

# ODDS RATIO

if crosstab_smoking.shape == (2, 2):

    a, b, c, d = crosstab_smoking.values.flatten()

    odds_ratio = (a * d) / (b * c)

    print("\nODDS RATIO:", odds_ratio)

    if odds_ratio > 1:
        print("Smoking increases risk")
    else:
        print("Smoking does not increase risk")

# 7. CORRELATION ANALYSIS

# ENCODE CATEGORICAL VARIABLES

encoded_df = df.copy()

mapping = {
    'Yes': 1,
    'No': 0,
    'High': 2,
    'Medium': 1,
    'Low': 0
}

for col in encoded_df.columns:

    if encoded_df[col].dtype == 'object':

```

```

        encoded_df[col] = encoded_df[col].map(mapping)

# CORRELATION MATRIX

corr = encoded_df.corr(numeric_only=True)

print("\nCORRELATION MATRIX")
print(corr)

sns.heatmap(
    corr,
    annot=True,
    cmap='coolwarm'
)

plt.title("Correlation Matrix")

plt.show()

# 8. LOGISTIC REGRESSION

# CONVERT TARGET VARIABLE TO NUMERIC

# Convert Yes/No → 1/0

df['LungDisease'] = df['LungDisease'].map({
    'Yes': 1,
    'No': 0
})

# DEFINE TARGET VARIABLE

y = df['LungDisease']

# DEFINE FEATURES

X = df.drop(columns=['LungDisease'])

# CONVERT CATEGORICAL VARIABLES

X = pd.get_dummies(X, drop_first=True)

# CONVERT TO NUMERIC

```

```

X = X.astype(float)

y = y.astype(int)

# ADD CONSTANT

X = sm.add_constant(X)

# CHECK DATA TYPES

print(X.dtypes)

# FIT LOGISTIC REGRESSION

log_model = sm.Logit(y, X).fit()

# SUMMARY

print(log_model.summary())

# ODDS RATIOS

odds_ratios = pd.DataFrame({
    'Variable': log_model.params.index,
    'Odds Ratio': np.exp(log_model.params.values)
})

print("\nODDS RATIOS")
print(odds_ratios)

```

FIRST 5 ROWS

	ID	Smoking	Age	Income	Pollution	LungDisease
0	1	Yes	36	Low	High	Yes
1	2	No	28	Low	Low	No
2	3	No	52	Medium	High	No
3	4	No	39	High	High	No
4	5	Yes	32	Low	Low	Yes

COLUMN NAMES

```
Index(['ID', 'Smoking', 'Age', 'Income', 'Pollution', 'LungDisease'],
      dtype='object')
```

DATA INFO

```
<class 'pandas.core.frame.DataFrame'>
```


RangeIndex: 500 entries, 0 to 499

Data columns (total 6 columns):

#	Column	Non-Null Count	Dtype
0	ID	500 non-null	int64
1	Smoking	500 non-null	object
2	Age	500 non-null	int64
3	Income	500 non-null	object
4	Pollution	500 non-null	object
5	LungDisease	500 non-null	object

dtypes: int64(2), object(4)


memory usage: 23.6+ KB

None

AGE SUMMARY

count	500.000000
mean	43.904000
std	14.604841
min	20.000000
25%	30.750000
50%	45.000000
75%	56.000000
max	69.000000

Name: Age, dtype: float64




SMOKING COUNTS

Smoking

No 293

Yes 207

Name: count, dtype: int64



INCOME GROUP COUNTS

Income

Low 204

Medium 197

High 99

Name: count, dtype: int64

MOST COMMON INCOME GROUP


Low

HIGH POLLUTION EXPOSURE: 70.40%

LUNG DISEASE PREVALENCE: 0.00%


CROSSTAB: SMOKING VS LUNG DISEASE

LungDisease	No	Yes
Smoking		
No	184	109
Yes	52	155



LUNG DISEASE BY INCOME GROUP (%)

LungDisease	No	Yes
Income		
High	51.515152	48.484848
Low	44.117647	55.882353
Medium	48.223350	51.776650



CHI-SQUARE TEST

Chi-square statistic: 67.5942649754021

P-value: 2.0085094601569646e-16

Significant association exists

POLLUTION VS LUNG DISEASE

Chi-square statistic: 33.84256825473482

P-value: 5.975692575712951e-09

Use Fisher's Exact Test when:

- Sample size is small
- Expected counts are low
- Data is 2x2

FISHER'S EXACT TEST

Odds Ratio: 5.031757233592096


P-value: 5.128610707567623e-17

ODDS RATIO: 5.031757233592096

Smoking increases risk

CORRELATION MATRIX

	ID	Smoking	Age	Income	Pollution	LungDisease
ID	1.000000	-0.062467	0.023478	0.042987	-0.044530	0.025980
Smoking	-0.062467	1.000000	0.052564	-0.029956	0.064685	0.371747
Age	0.023478	0.052564	1.000000	0.107329	0.040776	0.249694
Income	0.042987	-0.029956	0.107329	1.000000	-0.006312	-0.056439
Pollution	-0.044530	0.064685	0.040776	-0.006312	1.000000	0.264552
LungDisease	0.025980	0.371747	0.249694	-0.056439	0.264552	1.000000



const	float64
ID	float64
Age	float64
Smoking_Yes	float64

```

Income_Low      float64
Income_Medium   float64
Pollution_Low  float64
dtype: object
Optimization terminated successfully.
      Current function value: 0.550770
      Iterations 6

```

Logit Regression Results

```

=====
Dep. Variable:          LungDisease   No. Observations:          500
Model:                  Logit         Df Residuals:              493
Method:                 MLE          Df Model:                  6
Date:                  Sat, 16 May 2026   Pseudo R-squ.:            0.2036
Time:                  10:54:43         Log-Likelihood:           -275.38
converged:              True           LL-Null:                  -345.79
Covariance Type:       nonrobust       LLR p-value:              6.765e-28
=====

```

```

=
              coef      std err          z      P>|z|      [0.025
0.975]
-----
-
const          -2.4734      0.468      -5.281      0.000      -3.391
-1.555
ID              0.0011      0.001       1.471      0.141      -0.000
0.002
Age            0.0399      0.007       5.446      0.000       0.026
0.054
Smoking_Yes    1.7302      0.220       7.863      0.000       1.299
2.162
Income_Low     0.4561      0.285       1.598      0.110      -0.103
1.016
Income_Medium  0.2276      0.285       0.797      0.425      -0.332
0.787
Pollution_Low -1.3286      0.237      -5.597      0.000      -1.794
-0.863
=====
=

```

ODDS RATIOS

	Variable	Odds Ratio
0	const	0.084299
1	ID	1.001066
2	Age	1.040655
3	Smoking_Yes	5.642025
4	Income_Low	1.577917
5	Income_Medium	1.255540
6	Pollution_Low	0.264849

2. Suppose we want to determine if three different exercise programs impact weight loss differently. We recruit 90 people to participate in an experiment in which we randomly assign (using uniform distribution) 30 people to follow either program A, program B, or program C for one month. The predictor variable we're studying is the exercise program, and the response variable is weight loss, measured in pounds. Conduct a one-way ANOVA to determine if there is a statistically significant difference between the resulting weight loss from the three programs, check the model assumptions, and analyze treatment differences.

```
[11]: #Read Data
import numpy as np
import pandas as pd

np.random.seed(123)

# Generate uniform data
A = np.random.uniform(3, 7, 30)    # Program A
B = np.random.uniform(4, 8, 30)    # Program B
C = np.random.uniform(5, 9, 30)    # Program C

# Combine into DataFrame
data = pd.DataFrame({
    'weight_loss': np.concatenate([A, B, C]),
    'program': ['A']*30 + ['B']*30 + ['C']*30
})

print(data)
#Summary Data

data.groupby('program')['weight_loss'].describe()
#Normality Check
from scipy.stats import shapiro

for group in ['A', 'B', 'C']:
    stat, p = shapiro(data[data['program'] == group]['weight_loss'])
    print(f"Program {group}: p-value = {p:.4f}")

#Variance check
from scipy.stats import levene

stat, p = levene(A, B, C)
print("Levene p-value:", p)

#One way ANOVA
from scipy.stats import f_oneway

stat, p = f_oneway(A, B, C)
```



```

print("F-statistic:", stat)
print("p-value:", p)

#Visualization
import seaborn as sns
import matplotlib.pyplot as plt

sns.boxplot(x='program', y='weight_loss', data=data)
plt.title("Weight Loss by Program")
plt.show()


```

	weight_loss	program
0	5.785877	A
1	4.144557	A
2	3.907406	A
3	5.205259	A
4	5.877876	A
..	...	...
85	8.366680	C
86	6.429590	C
87	5.174366	C
88	6.219072	C
89	6.592743	C

```

[90 rows x 2 columns]
Program A: p-value = 0.7053
Program B: p-value = 0.1988
Program C: p-value = 0.1094
Levene p-value: 0.4801884167578272
F-statistic: 29.01563599838975
p-value: 2.2156719226973416e-10

```



Interpretation:

Result: The p-value < 0.05 suggests we reject the null hypothesis. Not all programs are equal. Post-hoc: Program B shows significantly higher weight loss than A or C.

3. Consider the Anemia Levels in Nigeria dataset: (<https://www.kaggle.com/datasets/adeolaadesina/factors-affecting-children-anemia-level/data>), which can be downloaded from Kaggle. The dataset comes from the 2018 Nigeria Demographic and Health Surveys (NDHS). It explores the impact of mothers' age and socioeconomic factors on anemia levels among children aged 0-59 months across Nigeria's 36 states and the Federal Capital Territory. Based on the above data, create a contingency table, apply chi-square test, formulate hypothesis, create a contribution diagram, and interpret the findings.

```
[12]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import chi2_contingency

#Load the dataset
df = pd.read_csv('children_anemia.csv')
```

```

#Data Cleaning
# We focus on 'Wealth index combined' and 'Anemia level.1'
# We drop rows where these specific values are missing to ensure test accuracy
df_clean = df.dropna(subset=['Wealth index combined', 'Anemia level.1'])
print(df_clean)

# 3. Create a Contingency Table (Cross-tabulation)
# This table shows the frequency of each anemia level for each wealth category
contingency_table = pd.crosstab(df_clean['Wealth index combined'],
                                df_clean['Anemia level.1'])

print("--- Observed Contingency Table ---")
print(contingency_table)
print("-" * 30)

# 4. Perform the Chi-Square Test of Independence
chi2, p, dof, expected = chi2_contingency(contingency_table)

print(f"Chi-Square Statistic: {chi2:.4f}")
print(f"Degrees of Freedom: {dof}")
print(f"P-value: {p:.4e}")

# Decision making at 5% significance level
alpha = 0.05
if p < alpha:
    print("\nResult: Significant (Reject H0)")
    print("Conclusion: There is a statistically significant relationship between_
    ↪Wealth Index and Anemia Level.")
else:
    print("\nResult: Not Significant (Fail to reject H0)")
    print("Conclusion: There is no significant relationship between Wealth Index_
    ↪and Anemia Level.")

# 6. Visualization (Stacked Bar Chart)
# Normalizing the data to show proportions (percentages) makes the relationship_
    ↪easier to see
plt.figure(figsize=(10, 6))
props = contingency_table.div(contingency_table.sum(axis=1), axis=0) * 100
props.plot(kind='bar', stacked=True, colormap='viridis')

plt.title('Percentage Distribution of Anemia Levels by Wealth Index',_
    ↪fontsize=14)
plt.xlabel('Wealth Index Category', fontsize=12)
plt.ylabel('Percentage (%)', fontsize=12)
plt.legend(title='Anemia Status', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.xticks(rotation=0)
plt.tight_layout()

```

	Age in 5-year groups	Type of place of residence \
3	25-29	Urban
5	30-34	Urban
6	35-39	Urban
9	20-24	Urban
12	25-29	Urban
...	...	...
33918	35-39	Rural
33919	35-39	Rural
33920	25-29	Rural
33921	25-29	Rural
33922	20-24	Rural

	Highest educational level	Wealth index combined \
3	Secondary	Richest
5	Higher	Richest
6	Secondary	Richest
9	Secondary	Richest
12	Higher	Richest
...	...	...
33918	Secondary	Richer
33919	Secondary	Richer
33920	No education	Richer
33921	Higher	Richer
33922	Secondary	Richer

	Births in last five years	Age of respondent at 1st birth \
3	1	25
5	1	30
6	2	32
9	1	19
12	1	24
...	...	...
33918	2	19
33919	2	19
33920	1	27
33921	1	22
33922	1	21

	Hemoglobin level adjusted for altitude and smoking (g/dl - 1 decimal) \
3	95.0
5	113.0
6	121.0
9	108.0
12	116.0
...	...
33918	120.0
33919	120.0

33920	120.0
33921	149.0
33922	123.0

	Anemia level \
3	Moderate
5	Mild
6	Not anemic
9	Moderate
12	Mild
...	...
33918	Not anemic
33919	Not anemic
33920	Not anemic
33921	Not anemic
33922	Not anemic

	Have mosquito bed net for sleeping (from household questionnaire) \
3	Yes
5	Yes
6	Yes
9	Yes
12	Yes
...	...
33918	Yes
33919	Yes
33920	Yes
33921	Yes
33922	Yes

	Smokes cigarettes	Current marital status \
3	No	Married
5	No	Married
6	No	Married
9	No	Married
12	No	Married
...	...	...
33918	No	Married
33919	No	Married
33920	No	Never in union
33921	No	Married
33922	No	Married

	Currently residing with husband/partner	When child put to breast \
3	Living with her	105.0
5	Living with her	NaN
6	Living with her	Immediately
9	Living with her	Immediately

12	Living with her	Days: 1
...	...	...
33918	Living with her	Immediately
33919	Living with her	NaN
33920	NaN	Hours: 1
33921	Living with her	Hours: 1
33922	Living with her	Immediately

	Had fever in last two weeks \
3	No
5	No
6	No
9	No
12	No
...	...
33918	No
33919	No
33920	No
33921	No
33922	No

	Hemoglobin level adjusted for altitude (g/dl - 1 decimal) \
3	114.0
5	119.0
6	102.0
9	113.0
12	109.0
...	...
33918	108.0
33919	120.0
33920	120.0
33921	119.0
33922	75.0

	Anemia level.1 Taking iron pills, sprinkles or syrup
3	Not anemic No
5	Not anemic No
6	Mild Yes
9	Not anemic Yes
12	Mild No
...	...
33918	Mild Yes
33919	Not anemic Yes
33920	Not anemic No
33921	Not anemic No
33922	Moderate Yes

[10182 rows x 17 columns]

--- Observed Contingency Table ---

Anemia level.1	Mild	Moderate	Not anemic	Severe
Wealth index combined				
Middle	596	865	729	59
Poorer	525	926	498	87
Poorest	522	983	418	119
Richer	606	750	737	48
Richest	505	403	797	9

Chi-Square Statistic: 530.3991

Degrees of Freedom: 12


P-value: 7.4510e-106

Result: Significant (Reject H0)

Conclusion: There is a statistically significant relationship between Wealth Index and Anemia Level.

<Figure size 1000x600 with 0 Axes>

Percentage Distribution of Anemia Levels by Wealth Index



Interpretation:

Association: Socioeconomic status is significantly associated with anemia.

Contribution: The “Poorest” group is significantly over-represented in severe anemia cases, while

the “Richest” group shows the lowest rates.

4. Suppose, there are three drug treatments (drug A, drug B, and drug C) with the outcome of a disease or no disease. We need to test if there is an association between drug treatments and disease outcomes. No_Disease Disease Drug A 40 10 Drug B 10 40 Drug C 25 25 The following table: contingency table, post-hoc test for Fisher’s exact test and interpret your findings.

```
[13]: import numpy as np
import pandas as pd
from scipy.stats import fisher_exact
from statsmodels.stats.multitest import multipletests

# 1. Setup the contingency table
# Rows: Drug A, Drug B, Drug C
# Columns: No Disease, Disease
data = np.array([
    [40, 10], # Drug A
    [10, 40], # Drug B
    [25, 25]  # Drug C
])

drugs = ['Drug A', 'Drug B', 'Drug C']

# 2. Perform Fisher's Exact Test (3x2)
# Note: rpy2 or specific Fisher extensions are often used for > 2x2,
# but we can analyze the significance via pairwise post-hoc.
print("--- Fisher's Exact Test Analysis ---")
# 3. Post-hoc Pairwise Comparisons
pairs = [(0, 1), (1, 2), (0, 2)]
p_values = []
results = []

for i, j in pairs:
    subset = data[[i, j], :]
    _, p_val = fisher_exact(subset)
    p_values.append(p_val)
    results.append(f"{drugs[i]} vs {drugs[j]}")

# 4. Bonferroni Correction
reject, pvals_corrected, _, _ = multipletests(p_values, alpha=0.05,
method='bonferroni')

# Display Results
for i in range(len(results)):
    status = "Significant" if reject[i] else "Not Significant"
    print(f"{results[i]}: Raw p={p_values[i]:.6f}, Corrected_
p={pvals_corrected[i]:.6f} ({status})")
```


--- Fisher's Exact Test Analysis ---

Drug A vs Drug B: Raw p=0.000000, Corrected p=0.000000 (Significant)

Drug B vs Drug C: Raw p=0.003052, Corrected p=0.009156 (Significant)

Drug A vs Drug C: Raw p=0.003052, Corrected p=0.009156 (Significant)

Interpretation:

Drug A Success Rate: 80% (Best treatment).

Drug B Success Rate: 20% (Least effective).

Drug C Success Rate: 50% (Neutral effect).

5. A researcher is interested in how variables, such as GRE (Graduate Record Exam scores), GPA (grade point average) and prestige of the undergraduate institution, effect admission into graduate school. The response variable, admit/don't admit, is a binary variable. The data set is taken from <https://stats.idre.ucla.edu/stat/data/binary.csv>. Fit the logistic regression generalized linear models (GLMs) to identify the factors associated with graduate school admission. Interpret the result.

```
[1]: import pandas as pd
import numpy as np
import statsmodels.formula.api as smf

# 1. Load Data
df = pd.read_csv("https://stats.idre.ucla.edu/stat/data/binary.csv")

# 2. Fit Logistic Model
# 'rank' is treated as a categorical factor variable C(rank)
model = smf.logit(formula="admit ~ gre + gpa + C(rank)", data=df).fit()

# 3. View Results
print(model.summary())

# 4. Calculate Odds Ratios (Exponentiated Coefficients)
odds_ratios = np.exp(model.params)
print("\nOdds Ratios:\n", odds_ratios)
```

Optimization terminated successfully.

Current function value: 0.573147

Iterations 6

Logit Regression Results

```
=====
Dep. Variable:          admit    No. Observations:          400
Model:                  Logit    Df Residuals:              394
Method:                 MLE      Df Model:                  5
Date:                  Sat, 16 May 2026    Pseudo R-squ.:        0.08292
Time:                  11:17:37    Log-Likelihood:       -229.26
converged:              True      LL-Null:              -249.99
Covariance Type:        nonrobust    LLR p-value:          7.578e-08
=====
```

	coef	std err	z	P> z	[0.025	0.975]
Intercept	-3.9900	1.140	-3.500	0.000	-6.224	-1.756
C(rank) [T.2]	-0.6754	0.316	-2.134	0.033	-1.296	-0.055
C(rank) [T.3]	-1.3402	0.345	-3.881	0.000	-2.017	-0.663
C(rank) [T.4]	-1.5515	0.418	-3.713	0.000	-2.370	-0.733
gre	0.0023	0.001	2.070	0.038	0.000	0.004
gpa	0.8040	0.332	2.423	0.015	0.154	1.454

Odds Ratios:

```
Intercept      0.018500
C(rank) [T.2]   0.508931
C(rank) [T.3]   0.261792
C(rank) [T.4]   0.211938
gre             1.002267
gpa             2.234545
dtype: float64
```

Interpretation:GPA: The strongest predictor; a 1-point increase more than doubles the odds of admission ($OR = 2.23$).Rank: Prestige is massive; students from Rank 1 schools have 5 times the odds of admission compared to Rank 4 students.

problem 6:

Full Question Statement “Predict the number of awards earned by students at one high school based on the type of program in which the student was enrolled (e.g., vocational, general, or academic) and the score on their final exam in math. The data set is taken from https://stats.idre.ucla.edu/stat/data/poisson_sim.csv. Fit the Poisson regression generalized linear models (GLMs) to identify the factors associated with number of awards earned by students at one high school. Interpret the result.

```
[ ]: import pandas as pd
import numpy as np
import statsmodels.api as sm
import statsmodels.formula.api as smf

# 1. Load the data
url = "https://stats.idre.ucla.edu/stat/data/poisson_sim.csv"
df = pd.read_csv(url)

# 2. Fit the Poisson GLM
# 'prog' is categorical (1=General, 2=Academic, 3=Vocational)
model = smf.glm(formula="num_awards ~ math + C(prog)",
                 data=df,
                 family=sm.families.Poisson()).fit()

# 3. Calculate Incidence Rate Ratios (IRR) for interpretation
irr = np.exp(model.params)
```

```
print(model.summary())
print("\nIncidence Rate Ratios (IRR):\n", irr)
```

Interpretation:

Overdispersion: The Negative Binomial is used because the variance of absences is much larger than the mean.

Findings: Higher math scores correlate with fewer absences. The Academic track has the best attendance records.

problem 7: School administrators study the attendance behavior (number of days absent) as a function of program type (e.g., vocational, general, or academic) and standardized math scores. The data set is taken from https://stats.idre.ucla.edu/stat/data/nb_data.dta . Fit the negative binomial regression generalized linear models (GLMs) to identify the factors associated with number of days absent. Interpret the result.

```
[27]: import pandas as pd
import numpy as np
import statsmodels.api as sm
import statsmodels.formula.api as smf

# Load data
url = "https://stats.idre.ucla.edu/stat/data/nb_data.dta"

df = pd.read_stata(url)

# Convert prog to string
df["prog"] = df["prog"].astype(str)

# Fit model
model = smf.glm(
    formula="daysabs ~ math + prog",
    data=df,
    family=sm.families.NegativeBinomial()
).fit()

# Summary
print(model.summary())

# IRR
irr = np.exp(model.params)

print("\nIncident Rate Ratios (IRR):")
print(irr)
```

Generalized Linear Model Regression Results

```
=====
Dep. Variable:          daysabs   No. Observations:          314
```

```

Model: GLM Df Residuals: 310
Model Family: NegativeBinomial Df Model: 3
Link Function: Log Scale: 1.0000
Method: IRLS Log-Likelihood: -865.68
Date: Sat, 16 May 2026 Deviance: 350.98
Time: 11:02:34 Pearson chi2: 331.
No. Iterations: 6 Pseudo R-squ. (CS): 0.1926
Covariance Type: nonrobust

```

	coef	std err	z	P> z	[0.025	0.975]
Intercept	2.6150	0.200	13.054	0.000	2.222	3.008
prog[T.2.0]	-0.4408	0.185	-2.379	0.017	-0.804	-0.078
prog[T.3.0]	-1.2786	0.203	-6.284	0.000	-1.677	-0.880
math	-0.0060	0.003	-2.358	0.018	-0.011	-0.001

Incident Rate Ratios (IRR):

```

Intercept    13.667374
prog[T.2.0]   0.643551
prog[T.3.0]   0.278418
math          0.994030

```

dtype: float64

```

C:\Users\Lenovo\anaconda3\Lib\site-
packages\statsmodels\genmod\families\family.py:1367: ValueWarning: Negative
binomial dispersion parameter alpha not set. Using default value alpha=1.0.
  warnings.warn("Negative binomial dispersion parameter alpha not "

```

Interpretation:

Performance: Higher math scores correlate with fewer absences.

Program: Vocational and Academic students have significantly better attendance than those in the General track.

problem 8: The state wildlife biologists want to model how many fish are being caught by fishermen at a state park. Visitors are asked how long they stayed, how many people were in the group, whether there were children in the group, and how many fish were caught. Some visitors do not fish, but there is no data on whether a person fished or not. Some visitors who did fish did not catch any fish, so there are excess zeros in the data because of the people that did not fish. The data set is taken from <https://stats.idre.ucla.edu/stat/data/fish.csv>. Fit the Zero-Inflated Poisson regression generalized linear model (GLMs) to identify the factors associated with the number of fish caught. Interpret the result.

```

[18]: import pandas as pd
import statsmodels.api as sm

# 1. Load the data
url = "https://stats.idre.ucla.edu/stat/data/fish.csv"

```

```

df = pd.read_csv(url)

# 2. Define variables
# Response: count (number of fish caught)
# Predictors for the Poisson part (count): child, camper, persons
# Predictors for the Logit part (zero-inflation): child, persons
y = df['count']
x_count = sm.add_constant(df[['child', 'camper', 'persons']])
x_zero = sm.add_constant(df[['child', 'persons']])

# 3. Fit the Zero-Inflated Poisson Model
model = sm.ZeroInflatedPoisson(y, x_count, exog_infl=x_zero, inflation='logit').
    ↪fit()

print(model.summary())

```

Optimization terminated successfully.

Current function value: 3.022141

Iterations: 23

Function evaluations: 25

Gradient evaluations: 25

ZeroInflatedPoisson Regression Results

```

=====
Dep. Variable:          count    No. Observations:          250
Model:                ZeroInflatedPoisson    Df Residuals:          246
Method:                  MLE    Df Model:                  3
Date:                   Sat, 16 May 2026    Pseudo R-squ.:          0.3296
Time:                   10:57:17    Log-Likelihood:         -755.54
converged:              True    LL-Null:               -1127.0
Covariance Type:        nonrobust    LLR p-value:            1.007e-160
=====

```

```

===

```

	coef	std err	z	P> z	[0.025
0.975]					

inflate_const	0.9813	0.414	2.370	0.018	0.170
1.793					
inflate_child	1.8227	0.315	5.780	0.000	1.205
2.441					
inflate_persons	-0.8454	0.190	-4.456	0.000	-1.217
-0.474					
const	-0.8452	0.166	-5.103	0.000	-1.170
-0.521					
child	-1.1472	0.091	-12.539	0.000	-1.327
-0.968					
camper	0.7597	0.094	8.096	0.000	0.576
0.944					

persons	0.8337	0.043	19.587	0.000	0.750
0.917					

=====

===

Interpretation:

Interpretation:Non-Fishers: Groups with children are significantly more likely to be non-fishers.Catch Rate: Campers catch 2.14 times more fish than non-campers.

problem 9: The state wildlife biologists want to model how many fish are being caught by fishermen at a state park. Visitors are asked how long they stayed, how many people were in the group, whether there were children in the group, and how many fish were caught. Some visitors do not fish, but there is no data on whether a person fished or not. Some visitors who did fish did not catch any fish, so there are excess zeros in the data because of the people that did not fish. The data set is taken from <https://stats.idre.ucla.edu/stat/data/fish.csv>. Fit the Zero-Inflated Binomial regression generalized linear model (GLMs) to identify the factors associated with the number of awards earned by students at one high school. Interpret the result.

```
[19]: import pandas as pd
import statsmodels.api as sm

# 1. Load the dataset from the URL provided in image_2a67c1.jpg
url = "https://stats.idre.ucla.edu/stat/data/fish.csv"
df = pd.read_csv(url)

# 2. Define the response and predictors
# count: Number of fish caught
# Predictors for count part: child, camper, persons
# Predictors for inflation part: child, persons
y = df['count']
x_count = sm.add_constant(df[['child', 'camper', 'persons']])
x_zero = sm.add_constant(df[['child', 'persons']])

# 3. Fit the Zero-Inflated Negative Binomial (ZINB) Model
# Note: ZINB is often preferred over ZIP if variance > mean
model = sm.ZeroInflatedPoisson(y, x_count, exog_infl=x_zero, inflation='logit').
    ↪fit()
print(model.summary())

# Calculate Odds Ratios (OR) and Incident Rate Ratios (IRR)
params = model.params
print("\n--- Odds Ratios (Non-Fisher Probability) ---")
print(np.exp(params[params.index.str.contains('inflate')]))

print("\n--- Incident Rate Ratios (Catch Rate) ---")
print(np.exp(params[~params.index.str.contains('inflate')]))
```

Optimization terminated successfully.

Current function value: 3.022141
Iterations: 23
Function evaluations: 25
Gradient evaluations: 25

ZeroInflatedPoisson Regression Results

```
=====
Dep. Variable:          count    No. Observations:          250
Model:          ZeroInflatedPoisson    Df Residuals:          246
Method:          MLE    Df Model:          3
Date:          Sat, 16 May 2026    Pseudo R-squ.:          0.3296
Time:          10:57:25    Log-Likelihood:          -755.54
converged:          True    LL-Null:          -1127.0
Covariance Type:          nonrobust    LLR p-value:          1.007e-160
=====
```

```
===
              coef      std err          z      P>|z|      [0.025
0.975]
-----
---
inflate_const      0.9813      0.414      2.370      0.018      0.170
1.793
inflate_child      1.8227      0.315      5.780      0.000      1.205
2.441
inflate_persons     -0.8454      0.190     -4.456      0.000     -1.217
-0.474
const              -0.8452      0.166     -5.103      0.000     -1.170
-0.521
child              -1.1472      0.091    -12.539      0.000     -1.327
-0.968
camper              0.7597      0.094      8.096      0.000      0.576
0.944
persons            0.8337      0.043     19.587      0.000      0.750
0.917
=====
===
```

--- Odds Ratios (Non-Fisher Probability) ---

```
inflate_const      2.667858
inflate_child      6.188700
inflate_persons     0.429402
dtype: float64
```

--- Incident Rate Ratios (Catch Rate) ---

```
const      0.429457
child      0.317512
camper     2.137712
persons    2.301896
dtype: float64
```

Interpretation: The “Zero” Factor: The presence of children is the primary reason groups don’t fish at all. Groups with kids are over 6 times more likely to be “non-fishers.”

The “Catch” Factor: For those who actually fish, campers and larger groups catch significantly more fish. However, even if they fish, having children in the group reduces the expected catch by about 68%.

Bottom Line: To catch the most fish, go in a large group of adult campers.

Problem 10: A study of length of hospital stay, in days, as a function of age, kind of health insurance and whether or not the patient died while in the hospital. Length of hospital stay is recorded as a minimum of at least one day. The data set is taken from <https://stats.idre.ucla.edu/stat/data/ztp.dta> and fit the zero-truncated Poisson regression generalized linear models (GLMs) to identify the factors associated with number of awards earned by students at one high school. Interpret the result.

```
[20]: import pandas as pd
import pyreadstat
import requests

# Download file
url = "https://stats.oarc.ucla.edu/stat/data/ztp.dta"

r = requests.get(url)

with open("ztp.dta", "wb") as f:
    f.write(r.content)

# Read stata file
df, meta = pyreadstat.read_dta("ztp.dta")

print(df.head())

# 2. Fit the Negative Binomial GLM
# Predictors: age, hmo (insurance type), died (binary outcome)
# Response: stay (number of days)
model = smf.glm('stay ~ age + hmo + died', data=df,
                family=sm.families.NegativeBinomial()).fit()

# 3. Model Summary
print(model.summary())

# 4. Calculate Incident Rate Ratios (IRR)
# IRR = exp(coefficient)
irr = np.exp(model.params)
print("\n--- Incident Rate Ratios (IRR) ---")
print(irr)
```

	stay	age	hmo	died
0	4	4	0	0.0

1	9	4	1	0.0
2	3	7	1	1.0
3	9	6	0	0.0
4	1	7	0	1.0

Generalized Linear Model Regression Results

```

=====
Dep. Variable:          stay    No. Observations:          1493
Model:                  GLM      Df Residuals:              1489
Model Family:      NegativeBinomial  Df Model:                3
Link Function:      Log         Scale:                  1.0000
Method:              IRLS       Log-Likelihood:        -4955.3
Date:                Sat, 16 May 2026    Deviance:              827.73
Time:                10:57:29    Pearson chi2:          938.
No. Iterations:      5          Pseudo R-squ. (CS):      0.01154
Covariance Type:      nonrobust
=====

```

	coef	std err	z	P> z	[0.025	0.975]
Intercept	2.4378	0.091	26.920	0.000	2.260	2.615
age	-0.0147	0.016	-0.894	0.371	-0.047	0.018
hmo	-0.1375	0.075	-1.845	0.065	-0.284	0.009
died	-0.2038	0.058	-3.505	0.000	-0.318	-0.090

--- Incident Rate Ratios (IRR) ---

```

Intercept    11.447342
age           0.985360
hmo           0.871547
died          0.815615

```

dtype: float64

C:\Users\Lenovo\anaconda3\Lib\site-

packages\statsmodels\genmod\families\family.py:1367: ValueWarning: Negative binomial dispersion parameter alpha not set. Using default value alpha=1.0.

warnings.warn("Negative binomial dispersion parameter alpha not "

[21]: !pip install pyreadstat

Requirement already satisfied: pyreadstat in c:\users\lenovo\anaconda3\lib\site-packages (1.3.4)

Requirement already satisfied: narwhals>=2.10.1 in

c:\users\lenovo\anaconda3\lib\site-packages (from pyreadstat) (2.21.0)

Requirement already satisfied: numpy in c:\users\lenovo\anaconda3\lib\site-packages (from pyreadstat) (2.3.5)

Interpretation:

Drivers: Insurance (HMO) and Death are the key predictors.

Findings: Patients who died or had HMO insurance had significantly shorter hospital stays than

survivors or privately insured patients.

Problem 11: A study of length of hospital stay, in days, as a function of age, kind of health insurance (HMO), and whether or not the patient died while in the hospital. Length of hospital stay is recorded as a minimum of at least one day. The data set is taken from <https://stats.idre.ucla.edu/stat/data/ztp.dta>. Fit the zero-truncated negative binomial regression generalized linear models (GLMs) to identify the factors associated with length of hospital stay. Interpret the result.

```
[22]: import pandas as pd
import numpy as np
import statsmodels.api as sm
import pyreadstat
import requests

from statsmodels.discrete.truncated_model import TruncatedLFNegativeBinomialP

# Download dataset
url = "https://stats.idre.ucla.edu/stat/data/ztp.dta"

r = requests.get(url)

with open("ztp.dta", "wb") as f:
    f.write(r.content)

# Read data
df, meta = pyreadstat.read_dta("ztp.dta")

print(df.head())

# Response variable
y = df["stay"]

# Predictor variables
X = df[["age", "hmo", "died"]]

# Add intercept
X = sm.add_constant(X)

# Fit Zero-Truncated Negative Binomial model
model = TruncatedLFNegativeBinomialP(y, X)

result = model.fit()

# Show results
print(result.summary())
```

```
      stay  age  hmo  died
0         4    4    0   0.0
```

```

1      9      4      1      0.0
2      3      7      1      1.0
3      9      6      0      0.0
4      1      7      0      1.0

```

Optimization terminated successfully.

Current function value: 3.185050

Iterations: 16

Function evaluations: 18

Gradient evaluations: 18

TruncatedLFNegativeBinomialP Regression Results

```

=====
=====
Dep. Variable:                stay    No. Observations:
1493
Model:                TruncatedLFNegativeBinomialP    Df Residuals:
1488
Method:                MLE    Df Model:
3
Date:                Sat, 16 May 2026    Pseudo R-squ.:
0.003263
Time:                10:57:33    Log-Likelihood:
-4755.3
converged:                True    LL-Null:
-4770.8
Covariance Type:                nonrobust    LLR p-value:
7.955e-07
=====
=====

```

	coef	std err	z	P> z	[0.025	0.975]
const	2.4083	0.072	33.457	0.000	2.267	2.549
age	-0.0157	0.013	-1.197	0.231	-0.041	0.010
hmo	-0.1471	0.059	-2.484	0.013	-0.263	-0.031
died	-0.2178	0.046	-4.718	0.000	-0.308	-0.127
alpha	0.5663	0.031	18.132	0.000	0.505	0.628

```

=====
=====

```

Interpretation: Significant Drivers: Mortality (Death) and HMO Insurance are the main factors that shorten hospital stays. Patients who pass away or have HMO coverage have stays that are 15% to 20% shorter than survivors or those with private insurance.

Non-Significant Factor: Age does not statistically impact the length of stay in this model.

Bottom Line: Hospital duration is driven by clinical outcomes and insurance type rather than the patient's age.

[]: